

LEC 05

CONCURRENT TCP SERVER

Bui Trong Tung, SoICT, HUST

1

Content

- I/O Models
- Concurrent TCP server: one child per client
- Concurrent TCP server: one thread per client

2

I/O MODELS

3

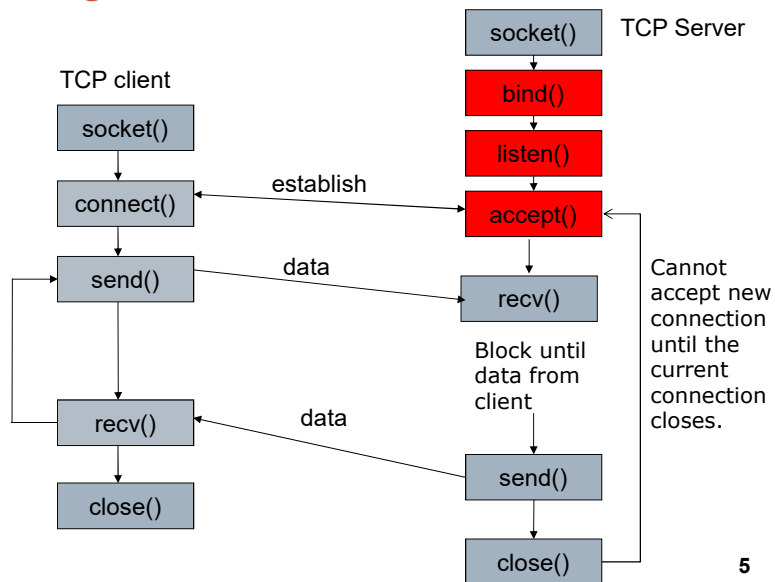
Review TCP Echo Server

```
while(1){
    //accept request
    connfd = accept(listenfd, (sockaddr *) & clientAddr,
                    &clientAddrLen);

    //receive message from client
    rcvBytes = recv(connfd, buff, BUFF_SIZE, 0);
    if(rcvBytes < 0){
        perror("Error :");
    }
    else{
        buff[rcvBytes] = '\0';
        printf("Receive from client: %s\n",buff);
        //Echo to client
        sendBytes = send(connfd, buff, strlen(buff), 0);
        if(sendBytes < 0)
            perror("Error: ",);
    }
    closesocket(connfd);
} //end while
```

4

Blocking I/O Model and TCP server

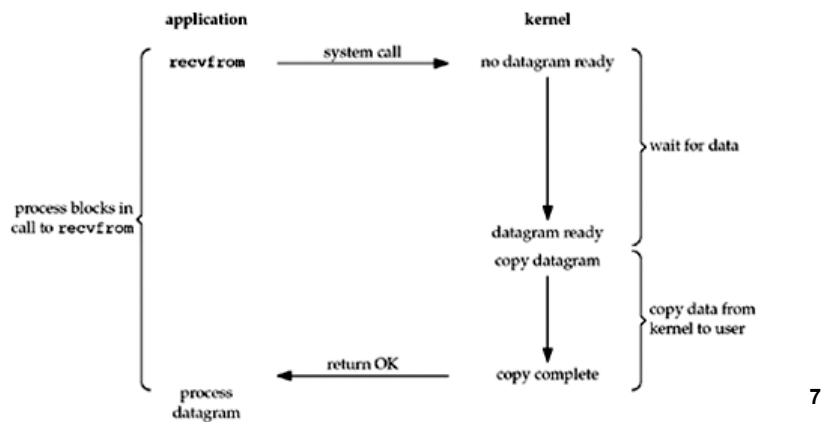


I/O Models

- blocking I/O
- nonblocking I/O
- I/O multiplexing (select and poll)
- signal driven I/O (SIGIO)
- asynchronous I/O (the POSIX aio_functions)

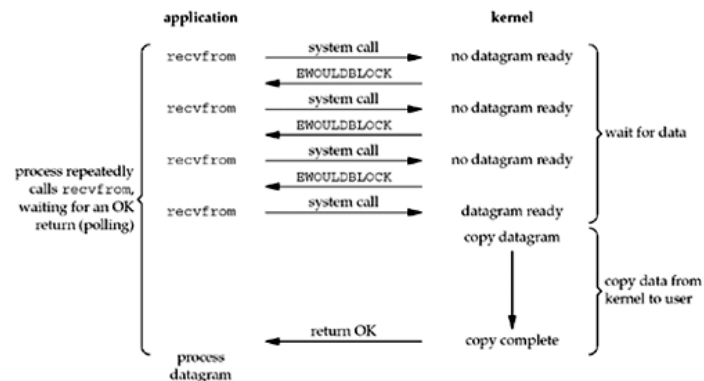
Blocking I/O Model

- Blocking I/O model: I/O function block process/thread until returning.
- `accept()`, `connect()`, `send()`, `recv()`, ...



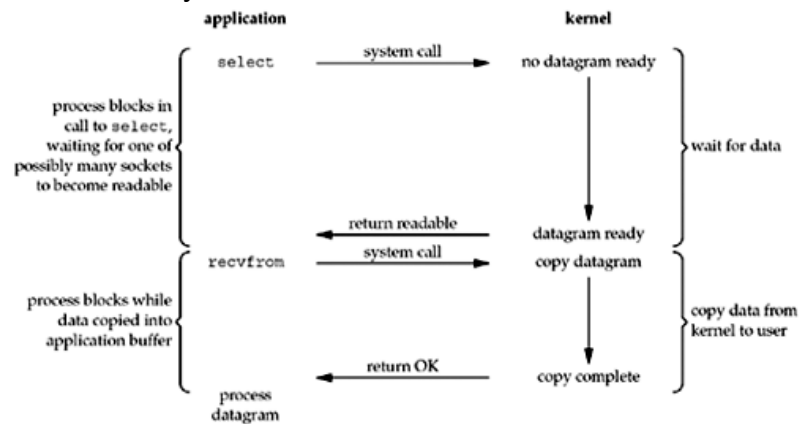
Non-blocking I/O Model

- Non-blocking I/O model: I/O function returns immediately
- If there is no data to return, so the kernel immediately returns an error of `EWOULDBLOCK` instead



I/O Multiplexing Model

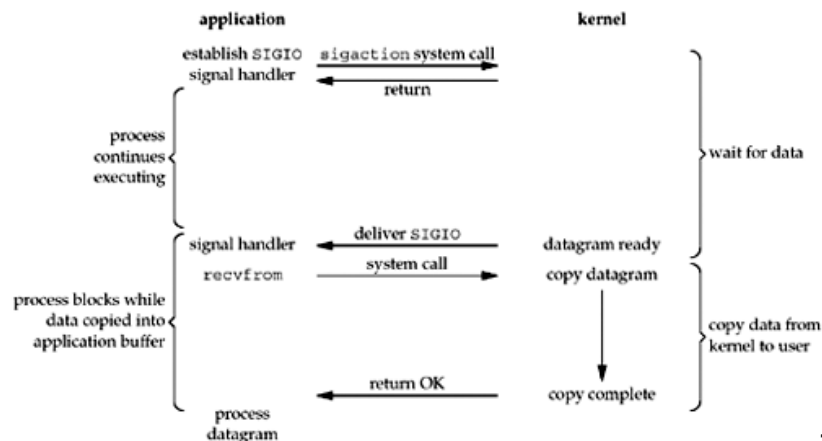
- With I/O multiplexing, we call `select` or `poll` and block in one of these two system calls, instead of blocking in the actual I/O system call



9

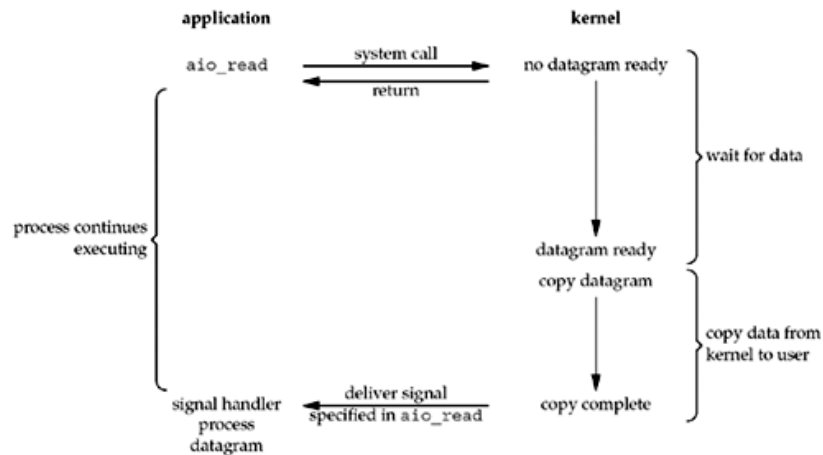
Signal-Driven I/O Model

- Use signals, telling the kernel to notify app with the `SIGIO` signal when the descriptor is ready



10

Asynchronous I/O Model



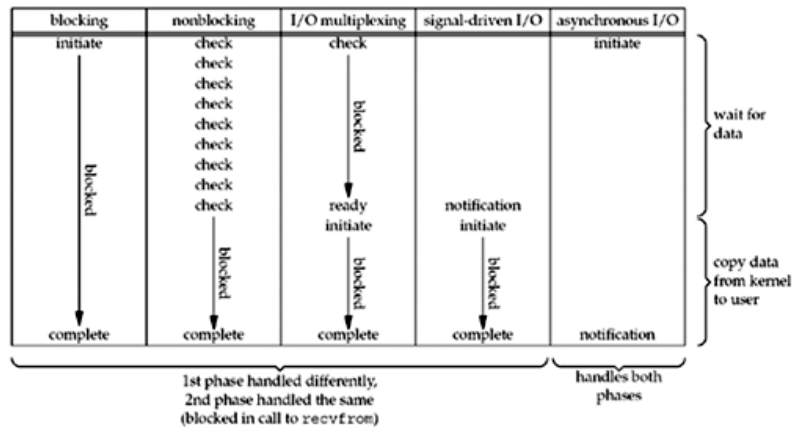
11

Asynchronous I/O Model (2)

- App calls `aio_read` (the POSIX asynchronous I/O functions begin with `aio_`)
- Pass the kernel
 - the descriptor
 - buffer pointer
 - buffer size (the same three arguments for `read`)
 - buffer offset (similar to `lseek`)
 - how to notify us when the entire operation is complete
- This system call returns immediately and our process is not blocked while waiting for the I/O to complete.

12

Comparison of the I/O Models



13

Iterating server

- Simple server
- But when a client request can take longer to service, we can't handle other clients
- Use a *concurrent server*
- One child per client: `fork()` a child process to handle each client
- One thread per client: create a thread to handle each client by using `pthread_create()`

14

MULTI-PROCESS SERVER

15


fork()

```
#include <unistd.h>
pid_t fork(void);
```

- Create a new process by copying itself.
- Returns twice:
 - Once in the calling process (called the parent) with a return value that is the process ID of the newly created process (the child).
 - Once in the child, with a return value of 0
- All descriptors open in the parent before the call to fork are shared with the child after fork returns


16

One child per client



17

One child per client



18

Use fork()

```
pid_t pid;
int listenfd, connfd;
//Step 1: Construct socket
//Step 2: Bind address to socket
//Step 3: Listen request from client

//Step 4: Communicate with client
while (1) {
    connfd = accept (listenfd, ... );
    if( (pid = fork()) == 0) { // process in child
        close(listenfd); // child closes listening socket
        doit(connfd);    // process the request
        close(connfd);   // done with this client
        exit(0);         // child terminates
    }
    close(connfd); // parent closes connected socket
}
```

19

Handling SIGCHLD Signals

- When a child process ends, it sends the **SIGCHLD** signal to the parent
 - Information about the child process is still maintained in “process table” in order to allow its parent to read the child exit status afterward.
- If we ignore the **SIGCHLD**, the child process will enter the zombie state
- We need to wait and handle **SIGCHLD** signal

20

Signaling

- A signal is a notification to a process that an event has occurred.
- Signals are sometimes called software interrupts.
- Signals usually occur asynchronously. By this we mean that a process doesn't know ahead of time exactly when a signal will occur.
- Signals can be sent
 - By one process to another process (or to itself)
 - By the kernel to a process

21

Signal (cont.)

- Typing certain key combinations at the controlling terminal of a running process causes the system to send it certain signals:
 - Ctrl-C sends an INT signal ("interrupt", SIGINT)
 - Ctrl-Z sends a TSTP signal ("terminal stop", SIGTSTP)
 - Ctrl-\ sends a QUIT signal (SIGQUIT)
- SIGHUP is sent to a process when its controlling terminal is closed (a hangup)
- SIGTERM is sent to a process to request its termination.
 - Unlike the SIGKILL signal, it can be caught and interpreted or ignored by the process.

22

Handling SIGCHLD Signals

- The purpose of the zombie state is to maintain information about the child for the parent to fetch at some later time.
- They take up space in the kernel and eventually we can run out of processes
- Whenever we *fork* children, we must *wait* for them to prevent them from becoming zombies → establish a signal handler to catch *SIGCHLD*, and within the handler, we call *wait*
- Establish the signal handler by adding the function call :
signal (SIGCHLD, handler);

23


wait() and *waitpid()*

```
#include <sys/wait.h>
pid_t wait (int *statloc);
pid_t waitpid (pid_t pid, int *statloc, int options);
```

- Wait for the status change of a process.
- Use to handle the terminated child
- Both return two values:
 - The return value of the function:
 - the process ID of the terminated child
 - 0 or -1 if error
 - The termination status of the child (an integer) is returned through the *statloc* pointer.

24

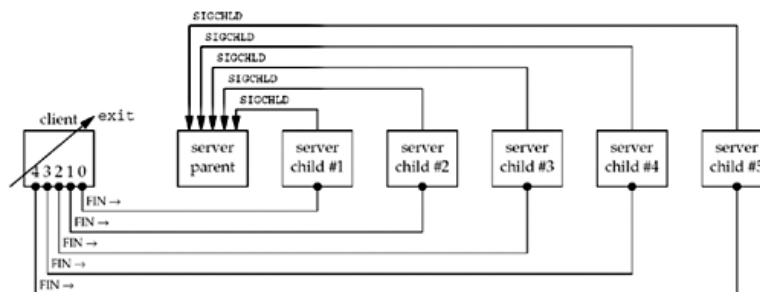
`wait()`



- Create 5 connections from a client to a forking server
- When the client terminates, all open descriptors are closed automatically by the kernel → five connections ended simultaneous

25

`waitpid()`



- Client terminates, closing all five connections, terminating all five children → four children are zombies
- It can happen when many users connect to a server
- → we have to use `waitpid()`

26

waitpid()

pid_t waitpid (pid_t pid, int *statloc, int options);

- pid < 0: wait for any child process whose process group ID is equal to the absolute value of *pid*.
- **pid = -1**: wait for any child process.
- pid = 0: wait for any child process whose process group ID is equal to that of the calling process
- pid > 0: wait for the child whose process ID is equal to the value of *pid*
- Without option WNOHANG, waitpid blocks until the status change
- With option WNOHANG, waitpid returns immediately
- Return
 - Pid of the child whose state has changed
 - with option WNOHANG, return 0 if the specified process has not changed status.

27

void sig_chld(int signo)

```
void sig_chld(int signo)
{
    pid_t pid;
    int stat;
    pid = waitpid(-1, &stat, WNOHANG );
    printf("child %d terminated\\", pid);
}
```

- **WNOHANG**: waitpid() does not block
- **while loop**: waitpid() repeatedly until there is no child process change status, i.e until waitpid returns 0.

28

Forking server

```
pid_t pid;
int listenfd, connfd;
//Step 1: Construct socket
//Step 2: Bind address to socket
//Step 3: Listen request from client

// wait for a child process to stop
signal(SIGCHLD, sig_chld);
//Step 4: Communicate with client
while (1) {
    connfd = accept (listenfd, ... );
    if ( pid = fork() == 0 ) { // process in child
        close(listenfd); // child closes listening socket
        doit(connfd);    // process the request
        close(connfd);   // done with this client
        exit(0);         // child terminates
    }
    close(connfd); // parent closes connected socket
}
```

29

Handling EINTR Errors

- When a process is blocked in a **slow system call** and the process catches a signal and the signal handler returns, the system call can return an error of EINTR.
- Slow system call: connect, accept, send, recv...
- Not all kernels automatically restart some interrupted system calls
- We must rewrite function to handle EINTR error

```
while(1)
    if ((connfd = accept (listenfd...)) < 0) {
        if (errno == EINTR)
            continue; /* back to for () */
        else
            perror ("Error: ");
    }
```

30

Other problems

- Connection abort before *accept* return
- Termination of server process
- Crashing of sever host
- Crashing and Reboot of server host

31

MULTI-THREAD SERVER

32

pthread_create()

```
#include <pthread.h>
int pthread_create(pthread_t *tid, const pthread_attr_t *attr,
                  void *(*routine) (void *), void *arg);
```

- Create a new thread
- Parameters:
 - [OUT] `tid`: points to ID of the new thread
 - [IN] `attr`: points to structure whose contents are used to determine attributes for the new thread
 - [IN] `routine`: the new thread starts execution by invoking `routine()`
 - [IN] `arg`: points to the argument is passed as the sole argument of `routine()`
- Return:
 - On success, returns 0
 - On error, returns an error number
- **Compile and link with `-pthread`**

33

pthread_create()

- By default, the new thread is joinable:
 - Not automatically cleaned up by GNU/Linux when it terminates
 - the thread's exit state hangs around in the system until another thread calls `pthread_join()` to obtain its return value
- Detached thread is cleaned up automatically when it terminates
 - Another thread may not obtain its return value
- Detach a thread: `int pthread_detach(pthread_t tid)`
 - On success, returns 0
 - On error, returns an error number

34

Multi-thread TCP Echo Server

```
pthread_t tid;
int listenfd, *connfd;
//Step 1: Construct socket
//Step 2: Bind address to socket
//Step 3: Listen request from client

//Step 4: Communicate with client
while (1) {
    connfd = malloc(sizeof(int));
    *connfd = accept (listenfd, ... );
    pthread_create(&tid, NULL, &client_handler, connfd);
}

close(listenfd);
return 0;
```

35

Multi-thread TCP Echo Server(cont.)

```
void *client_handler(void *arg){
    int connfd;
    int sendBytes, rcvBytes;
    char buff[BUFF_SIZE + 1];

    pthread_detach(pthread_self());
    *connfd = *((int *) arg);
    while(1){
        rcvBytes = recv(connfd, buff, BUFF_SIZE, 0);
        if (rcvBytes <= 0)
            break;

        sendBytes = send(connfd, buff, rcvBytes, 0);
        if (sendBytes <= 0)
            break;
    }
    close(connfd);
}
```

36

Synchronize threads

- Since multiple threads can be running concurrently, accessing the shared variables:
 - The order of the accessing shared memory is unpredictable, so
 - The processing flow of the thread may be uncontrollable, and/or
 - The process crash
- Synchronize threads so that only one thread can access shared memory:
 - Inter-lock
 - Semaphore
 - Mutex

37

Mutex

```
#include <pthread.h>
int pthread_mutex_lock(pthread_mutex_t * mptr);
int pthread_mutex_unlock(pthread_mutex_t * mptr);
```

- The thread can access the shared variable only when it hold the mutex
- pthread_mutex_lock(): lock a mutex
- pthread_mutex_unlock(): unlock a mutex
- If the thread try to lock a mutex that is already locked by some other thread, it is blocked until the mutex is unlocked.

```
void *routine(void *arg) {
    //...
    pthread_mutex_lock(mptr);
    // access shared memory
    pthread_mutex_unlock(mptr);
    //...
}
```

38

fork() VS pthread_create()

fork()

- Heavy-weight
- Passing information from the parent to the child before the fork is easy
- Returning information from the child to the parent takes more work
- Needn't synchronize processes
- Greater isolation between the parent and the child

pthread_create()

- Light-weight
- Passing information from a thread to the others is easy
- Don't need signal-driven processing when the threads ends.
- May synchronize threads
- If a thread crashes, process may crash

39